

# On Optimistic Methods for Concurrency Control

AUTHORS: H.T. KUNG and JOHN T. ROBINSON

---

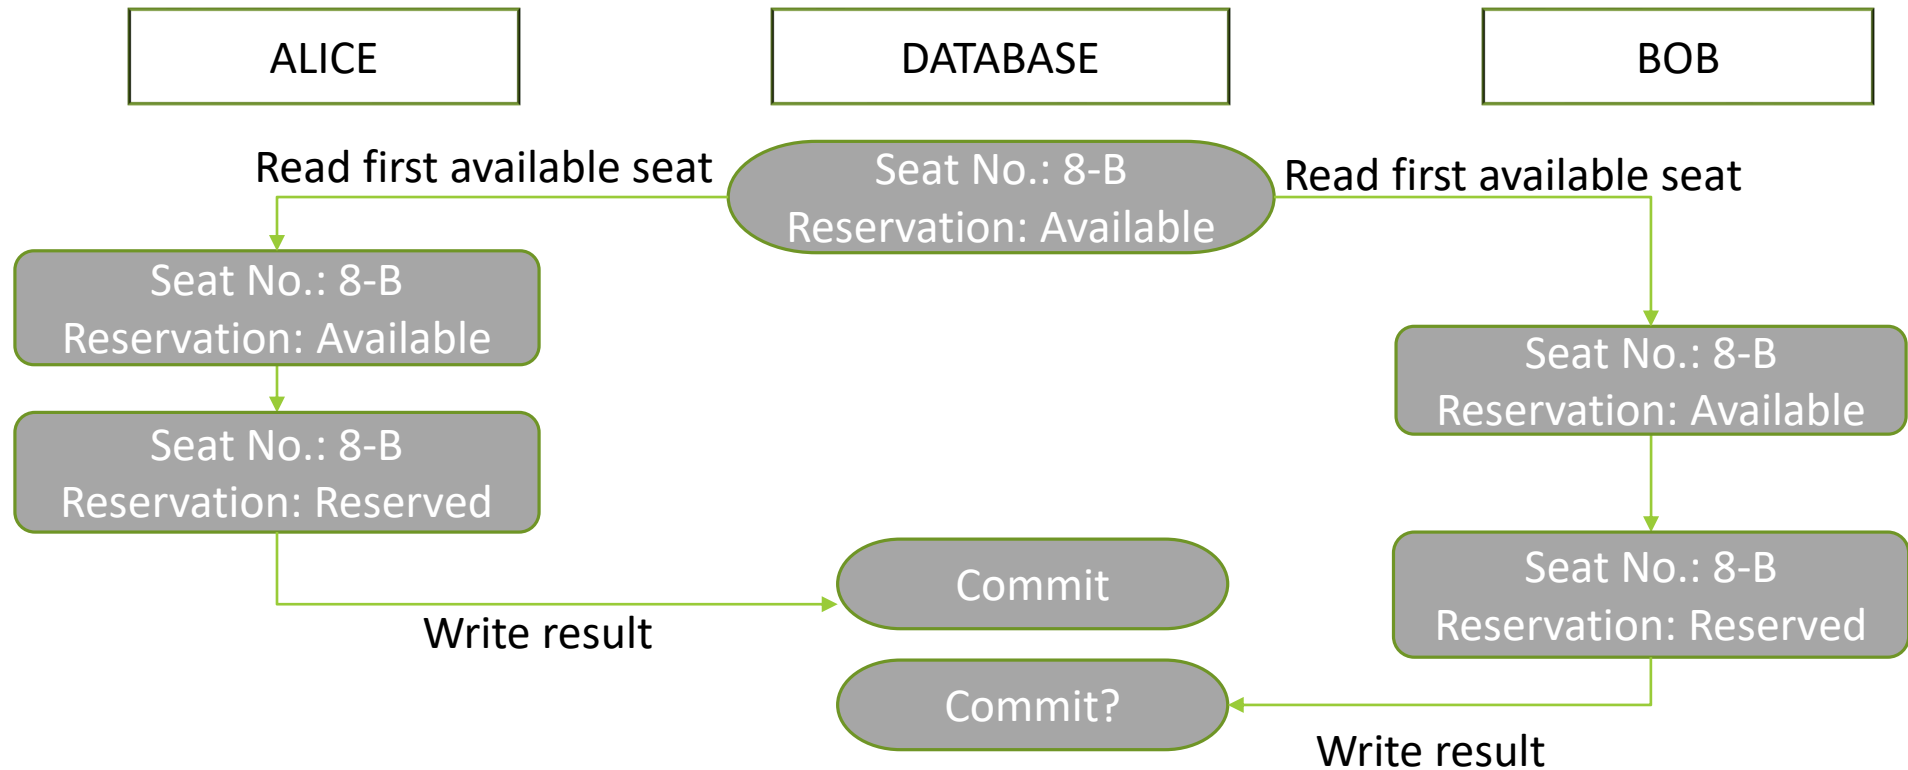
PRESENTATION BY:  
DHIVYA SIVARAMAKRISHNAN

# Introduction

---


- ❑ What is a Transaction?
- ❑ What is concurrent access?
- ❑ Why is it so desirable?
- ❑ Preserving database integrity while allowing concurrency

# Lost Update



# Locking Insights

- ❑ Alice and Bob access the database concurrently
- ❑ Locking - exclusive access to the resource
- ❑ Other attempts will be invalidated
- ❑ No process will act upon obsolete or work-in-progress information




# Trade-offs with Locking

---

- ❑ Lock maintenance overheads
- ❑ Impacts on concurrency, especially for aborted transactions
- ❑ Ensuring availability of congested nodes
- ❑ Secondary memory swaps on locked resources
- ❑ Deadlocks can occur

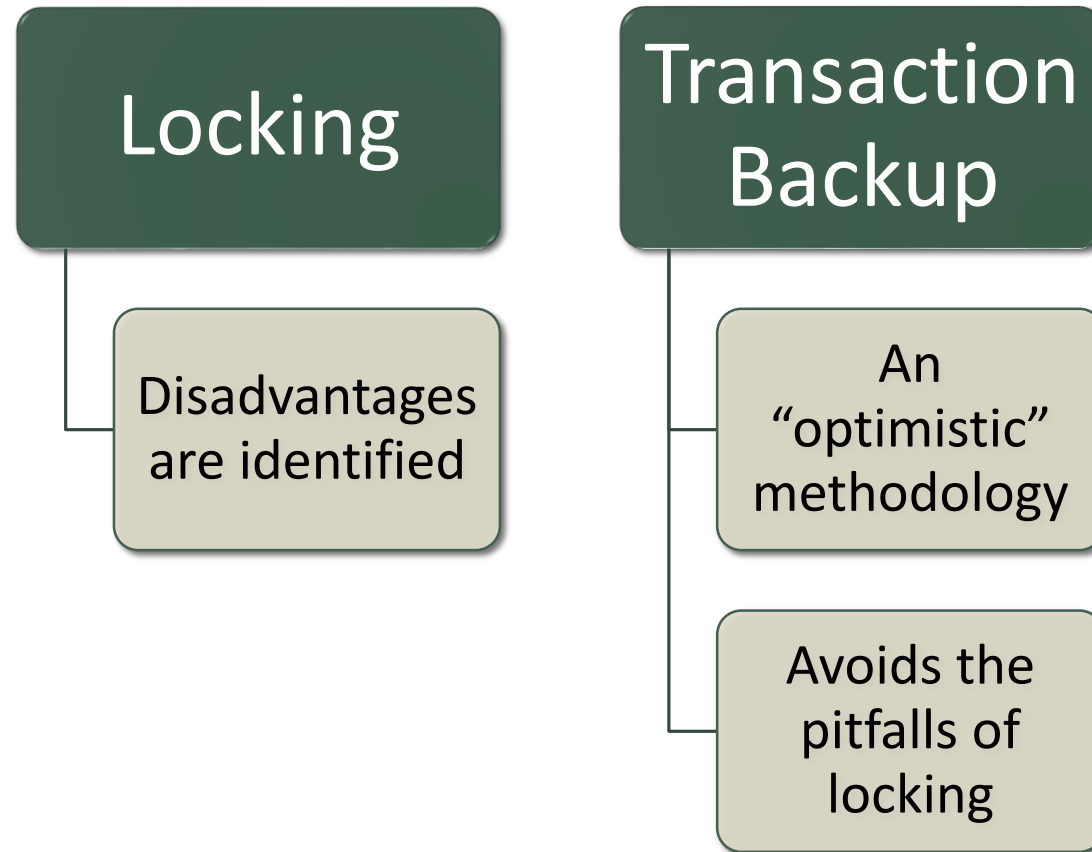
# Deadlocking Explained

---



# Focus of this Paper

---



# Foundation for “Optimism”

---

*Locking may be necessary only in the worst case*

General cases:

- ☐ Very high number of total resources compared to those being accessed
- ☐ Probability of modifying a congested resource is less
- ☐ Access conflicts will not happen among transactions



# Three Phases

---

Read

- Unrestricted - not a “modify” and cannot affect database integrity

Validation

- Determines if transaction causes any loss of integrity

Write

- Stringent restrictions. Writes happen only if validation succeeds

# What to Know

---

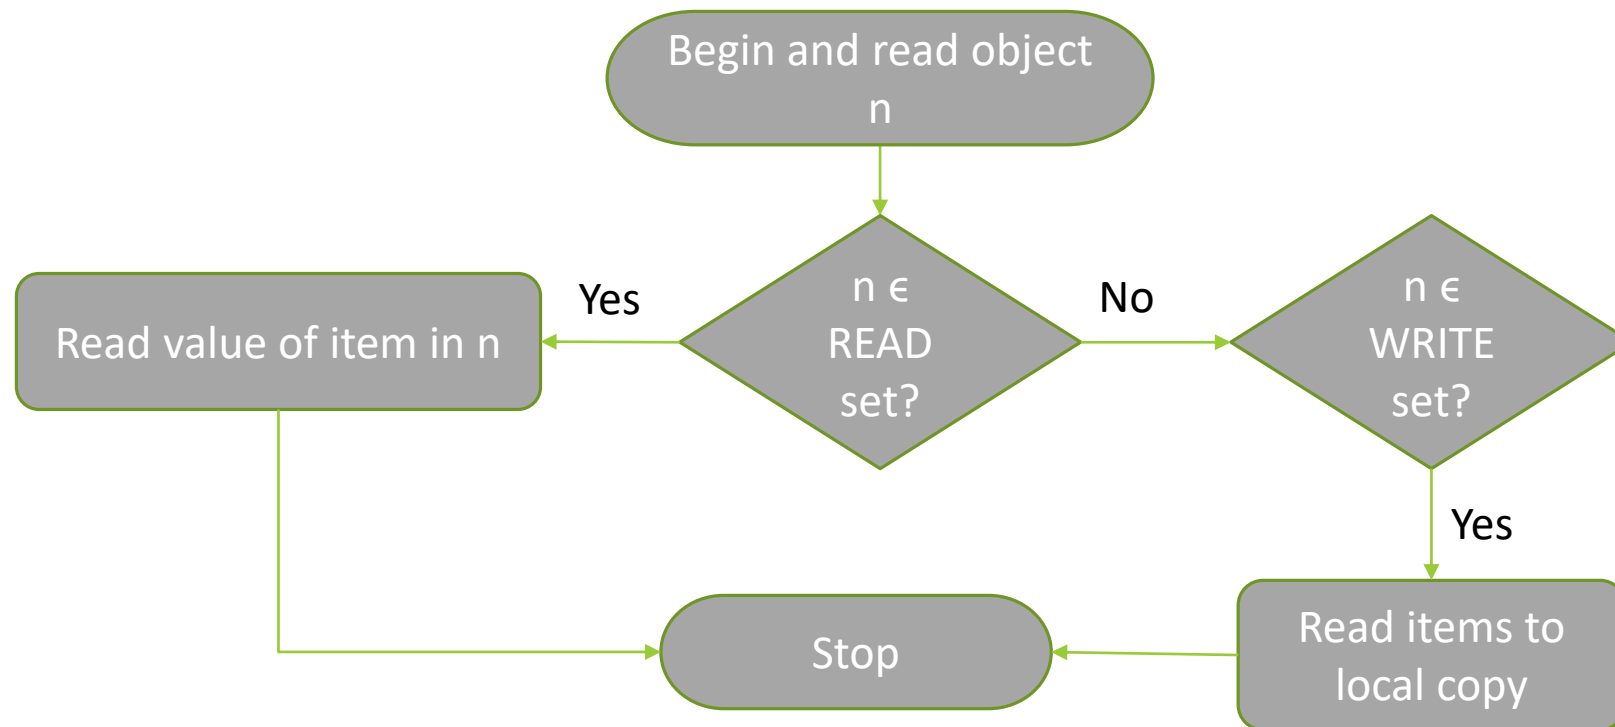
- ❑ A set of homogeneous objects of type A
- ❑ Concurrency control mechanism maintains OBJECT NAMES used by every transaction
- ❑ Assumed to be an empty set at the very beginning
- ❑ Every transaction has two copies of objects used – “read” set, “write” set

CREATE	Create a new object and return its name
READ	Read an item of an object
WRITE	Write a value to an item of an object
DELETE	Delete an object
COPY, EXCHANGE	Create copies of an object, swap values

# READ

# WRITE

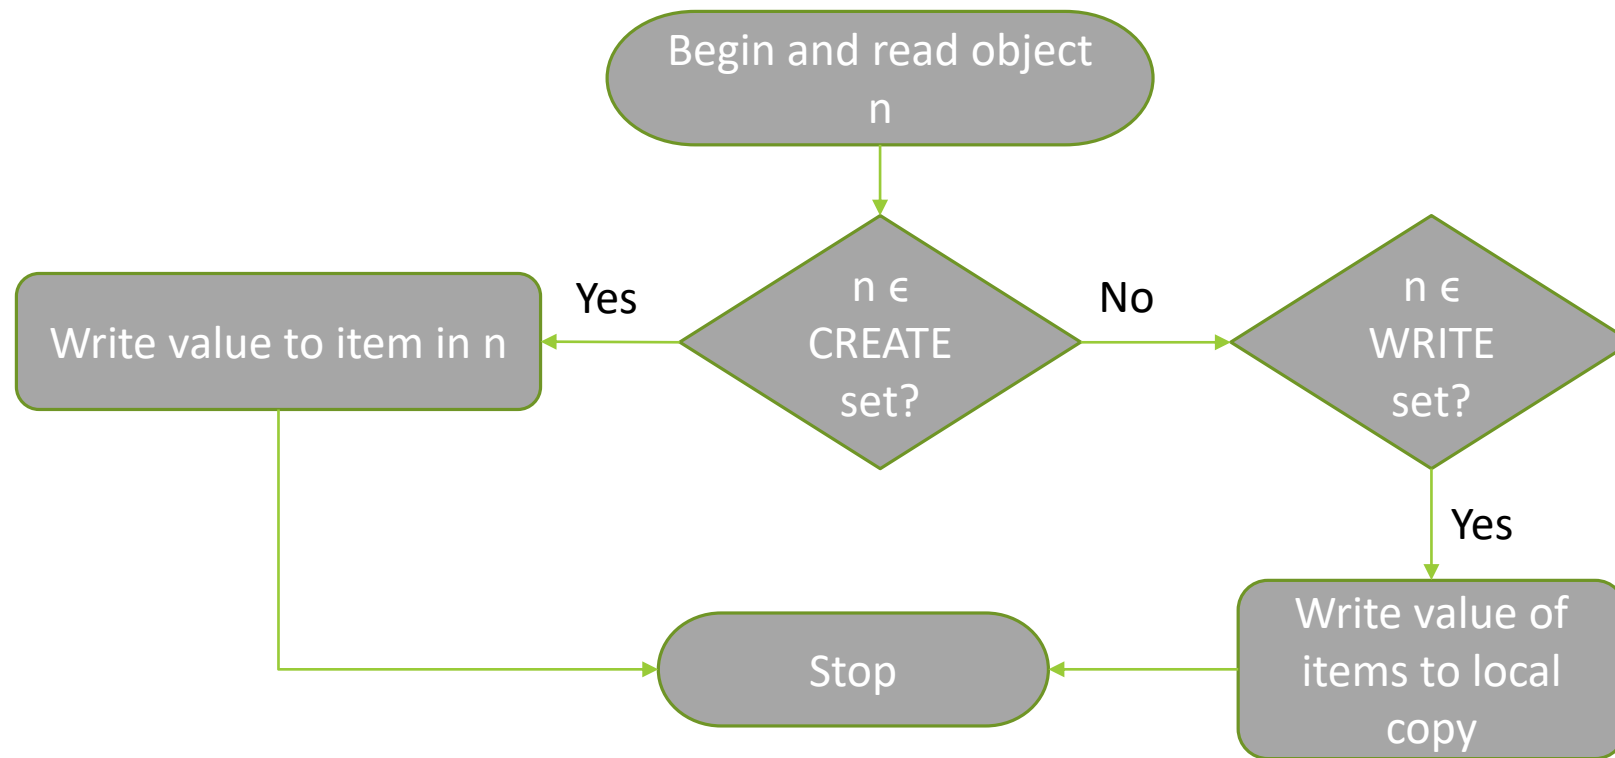
# VALIDATE



READ

WRITE

VALIDATE



# Integrity Preservation

---


- ❑ No root node can be created without writing new pointers to access it
- ❑ Root node deletions must clean up dangling pointers
- ❑ At transaction completion –
  - ❑ Created nodes become accessible
  - ❑ Deleted nodes become inaccessible
- ❑ Cleanup also happens after a transaction is aborted
- ❑ At the end of READ, all changes to 'n' are known

# READ

# WRITE

# VALIDATE

- ❑ Every transaction aims to preserve integrity of this shared data structure, D
- ❑ Check if D has been updated by any other transaction since the start of T1
- ❑ How do we verify the correctness of this concurrent execution?



# Serial Equivalence

---

- ❑ 'n' transactions concurrently access a database resource
- ❑ Two instances of transaction interleaving

Complete Schedule	
T <sub>1</sub>	T <sub>2</sub>
R(A)	
W(A)	
Commit	
	R(B)
	W(B)
	Abort

Complete Schedule	
T <sub>1</sub>	T <sub>2</sub>
R(A)	
	R(B)
W(A)	
	W(B)
Commit	
	Abort

- ❑ Same effect on database as if all the transactions ran one after the other



# Why is Serial Equivalence important?

---

- ❑ An easy way to validate that every transaction preserves integrity
- ❑ Easier to verify serial equivalence than check integrity after every interleaving of concurrent transactions
- ❑ Preserves the basic property for consistency – every transaction is atomic in nature
- ❑ Any amount of interleaving is possible, but the end result is the same – a consistent state

# Validating Serial Equivalence

---

- ❑ PROBLEM: Prove that database state remains same after any interleaving
- ❑ Find a permutation such that serial equivalence holds
- ❑ Assign transaction numbers  $t(\text{tname})$

$$t(i) < t(j)$$

# Transaction Numbers

---

Each transaction given a unique number

- Indicates its position in time

Number assigned through counters

- End of READ

Transactions that complete WRITE

- Number retained

Aborted transactions

- Number recycled

# Three Validation Conditions

---




Fig. 2. Possible interleaving of two transactions.

# Serial Validation

---

- ❑ First of a family of concurrency controls
- ❑ Utilizes validation conditions 1 and 2 - sequential WRITES
- ❑ Record “read” and “write” sets to local copy
- ❑ Tid, validation and subsequent write are all in a critical section

```

thegin = (
    create set := empty;
    read set := empty;
    write set := empty;
    delete set := empty;
    start tn := tnc)

tend = (
    (finish tn := tnc;
    valid := true;
    for t from start tn + 1 to finish tn do
        if (write set of transaction with transaction number t intersects read set)
            then valid := false;
    if valid
        then ((write phase); tnc := tnc + 1; tn := tnc));
    if valid
        then (cleanup)
        else (backup)).

```

# Parallel Validation

---

- ❑ Another family of concurrency control
- ❑ Uses all three validation conditions
- ❑ Multiple transactions may be in the validation phase at once
- ❑ Provides optimization similar to Serial Validation

# Parallel Validation - Procedure

---


- ❑ Save active transactions – finished READ
- ❑ Validate against conditions 1 and 2
- ❑ Validate against 3 for all transactions in “active” set
- ❑ If no conflicts, remove self from active and assign T(id)
- ❑ Else, abort



# A Comparison

## NO CONFLICTS

### Optimistic Concurrency Control




Allow computation to proceed without blocking.

Kung & Robinson. *On optimistic methods for concurrency control.*

## CONFLICT

### Optimistic Concurrency Control



Take a compensating action.

Kung & Robinson. *On optimistic methods for concurrency control.*

# A Quick Recap


---

## READ

- Create local copies

## WRITE

- Make local copies global



## VALIDATION

- Ensure database consistency

# Merits

---

- ❑ Locking overheads avoided - good access throughput
- ❑ Conflicts are assessed pretty early - at the end of READ
- ❑ Maximized parallelism
- ❑ Cost of rollbacks is lesser than deadlock resolution cost
- ❑ Negligible concurrency control overhead if more READs

# Major Demerits

---

- ❑ Relies solely on the belief that the likelihood of two transactions conflicting is low
- ❑ Conflicting transactions need to be aborted and restarted
- ❑ Too much redundancy if many transactions are aborted
- ❑ With heavy concurrency, heavy load and failure probabilities
- ❑ Starvation when same transactions are aborted

# Real-Time Users of OCC

---



# Conclusion

---

- ❑ Two branches of concurrency control
- ❑ Locking => resource-waiting
- ❑ Optimistic methods => all transactions to proceed and conflicting ones are aborted
- ❑ How to choose –
  - ❑ Locking – when chances of users updating same objects at once are high
  - ❑ Optimistic – if resources are many but transactions are fewer; more READs
- ❑ Unified goals – more throughput, less turnaround time



# References (In Order of Slides)

---

- 4 - <http://sharetraveler.com/get-good-seat-economy-class/>
- 6 - <http://csunplugged.org/routing-and-deadlock/>
- 11 - <https://slidehunter.com/powerpoint-templates/setting-smart-objectives-powerpoint-template/>
- 15 - [https://www.ibm.com/support/knowledgecenter/SSPHQG\\_7.2.0/com.ibm.powerha.concepts/ha\\_concepts\\_8node\\_takeover.htm](https://www.ibm.com/support/knowledgecenter/SSPHQG_7.2.0/com.ibm.powerha.concepts/ha_concepts_8node_takeover.htm)
- 16 - <https://gradeup.co/transactions-and-concurrency-control-i-4c5d9b27-c5a7-11e5-bcc4-bc86a005f7ba>
- 20, 22 - H. T. Kung, John T. Robinson, 'On Optimistic Methods for Concurrency Control', June 1981
- 25 - Concurrency Control for Parallel Machine Learning - <http://www.slideshare.net/jeykottalam/concurrency-control-for-machine-learning>
- 29 - <https://en.wikipedia.org>
- 31 - Google Images for "Questions"