

Eddies: Continuously Adaptive Query Processing

Ron Avnur

Joseph M. Hellerstein

Presented by: Samaneh Khakshour

CMPT 843

February 2017

Outline

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

❖ Motivation

❖ What are traditional query optimizers

❖ What are the problems with these optimizers

❖ What is eddy

❖ How to reorder operators

❖ Routing tuples in eddy

❖ Experimental results

❖ Summary

Static Query Processing

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

❖ Traditional query processing scheme:

❖ Query optimization

❖ Static query execution

❖ This scheme is not good for widely-distributed resources and massively parallel database systems.

Why not Static Query Processing

Motivation

- Traditional query optimizers

Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Large-scale systems should function robustly in an unpredictable environment with:
 - ❖ Hardware and workload complexity
 - ❖ Data complexity
 - ❖ User Interface complexity
- Query processing should be adaptive
 - Allowing the system to adapt dynamically to fluctuations in:
 - Computing resources
 - Data characteristics
 - User preferences

Run- Time Fluctuations

Motivation

- Traditional query optimizers
- **Problems**

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Three properties can vary during query processing:
 - ❖ Cost of Operators
 - ❖ Operator Selectivity
 - ❖ Rate tuple arrive from input
- ❖ Example on run-time variation in selectivity:
- ❖ Consider an employee table clustered by ascending age, and a selection “salary>10000”:
 - ❖ Initially the predicate “salary>10000” will be very selective
 - ❖ Selectivity rate will change as the older employees are scanned.

Eddies

Motivation


- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Two Challenges with Eddies:
 - ❖ How can we reorder operators (Reorderability of plan)
 - ❖ How should we route tuples?



Reorderability of Plans

Motivation

- Traditional query optimizers
- Problems

What is eddy

- **Reordering operators**
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Reoptimizing a query execution pipeline on the fly requires significant care in maintaining query execution state.
- ❖ Some things that Eddies must consider:
 - ❖ Adaptive or non existent synchronizations barriers
 - ❖ Frequent moment of symmetry

Synchronization Barriers

Motivation


- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ One Operation hinder the speed of other operations
- ❖ Example:
- ❖ The process of **fasthi** is postponed for a long time while consuming many tuples from **slowlow**.



- ❖ Desirable to minimize the number of synchronization barriers

Moments of Symmetry

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ When query is executed to a point that optimizer can change the query plan without affecting the way the query plans predicates are performed.

- ❖ Example:

- ❖ Consider a nested-loop join

- ❖ Outer relation R, and inner relation S

- ❖ Reordering the input (re-optimization):

- ❖ Can happen when S is completely scanned


- ❖ $R \bowtie S = S \bowtie R$

- ❖ Combination of

- ❖ Commutativity of operators

- ❖ Moment of symmetry

- Reordering of a plan



Join Algorithms and Reordering

Motivation

- Traditional query optimizers
- Problems

What is eddy

- **Reordering operators**
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Join ordering constrains:
 - ❖ Unindexed join input is ordered before the indexed input
 - ❖ Preserving the ordered inputs
 - ❖ Some join algorithm work only for equijoins
- ❖ In order for Eddies to be most effective:
 - ❖ Frequent moment of symmetry
 - ❖ Adaptive or non existent barriers
 - ❖ Minimal ordering constrains
- ❖ Ripple joins are having frequent moments of symmetry and attractive adaptivity.

Ripple Join


Motivation

- Traditional query optimizers
- Problems

What is eddy

- **Reordering operators**
- Eddies
- Routing tuples
- Experimental results

Summary



- ❖ Get tuples from each relation
- ❖ Compare them with tuples seen until now

Ripple Join




Motivation

- Traditional query optimizers
- Problems

What is eddy

- **Reordering operators**
- Eddies
- Routing tuples
- Experimental results

Summary



- ❖ Ripple joins:
 - ❖ Have moments of symmetry at each corner
 - ❖ Are designed to allow changing rates for each input
 - ❖ Offer attractive adaptivity feature

Eddy

Motivation


- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- **Eddies**
- Routing tuples
- Experimental results

Summary

- ❖ River: A shared nothing parallel query processing framework.
- ❖ Eddy is implemented via a module in River, containing:
 - ❖ arbitrary number of input relations
 - ❖ A number of participating unary and binary modules
 - ❖ A single output relation
- ❖ Eddy was designed to dynamically reoptimize queries.
- ❖ Eddy maintains a fixed size buffer of tuples that need to be processed



Eddy

Motivation


- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- **Eddies**
- Routing tuples
- Experimental results

Summary

- ❖ A tuple in eddy is associated with a descriptor:
 - ❖ A vector of Ready bits and Done bits
- ❖ The eddy forward the tuple to operators that have Ready bits turned on
- ❖ After an operator is processed, its Done bits are set
 - ❖ If all Done bits are set, tuple will be sent to output
 - ❖ Otherwise, sent to other operators



Routing Tuples in Eddies

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- **Routing tuples**
- Experimental results

Summary

- ❖ Eddy module:
 - ❖ Directs the flow of tuples from inputs through the various operators to the output
 - ❖ Providing the flexibility to allow each tuple to be routed individually through the operators
- ❖ Routing policy used in eddy determines the efficiency of the system
- ❖ In this paper multiple different ways of routing is studied.

Naïve Eddy

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- **Routing tuples**
- Experimental results

Summary

- ❖ Eddy's buffer is implemented as priority queue
- ❖ When a tuple enters buffer, its priority is set low
- ❖ After it is processed by an operation in Eddy and returned to buffer, its priority is set to high
- ❖ This ensure that tuples do not get clogged in the Eddy
- ❖ Eddy's queue size is fixed: back-pressure
 - ❖ Production along the input to any edge is limited by the rate of consumption at the output
 - ❖ Tuples are routed to low cost operators first
 - ❖ Consumption rate of low cost operator is higher than that of high cost operator.
- ❖ Cost aware policy
- ❖ Selectivity unaware policy

Lottery Scheme

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- **Routing tuples**
- Experimental results

Summary

- ❖ Gives priority to operators with low cost and low selectivity.
- ❖ Each time the eddy gives a tuple to an operator, it credits the operator one ticket.
 - Favor low cost (Track consumption)
- ❖ Each time the operator returns a tuple to eddy, one ticket is debited
 - from eddy's running count for that operator
 - Favor low selectivity (Track production)
- ❖ Operators chance of receiving a tuple corresponds to its count of tickets.
- ❖ Eddy can track an ordering of operators that give good overall efficiency.



Window Scheme

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- **Routing tuples**
- Experimental results

Summary

- ❖ Problem with lottery scheme is that it uses too much of past information:
 - ❖ An operator gained a lot of tickets initially but then became slow
- ❖ Window scheme: the lottery scheme is modified such that the lottery only looks at tickets gained by an operator in a fix window
 - ❖ Keep track of two types of tickets:
 - ❖ Banked tickets: Used when running the lottery
 - ❖ Escrow tickets: Used to measure efficiency during the window
 - ❖ At the end of the window:
 - ❖ Banked tickets=Escrow tickets
 - ❖ Escrow tickets=0
- ❖ Ensures operators re-ensure themselves each window



Experimental Results

Motivation

- Traditional query optimizers
- Problems

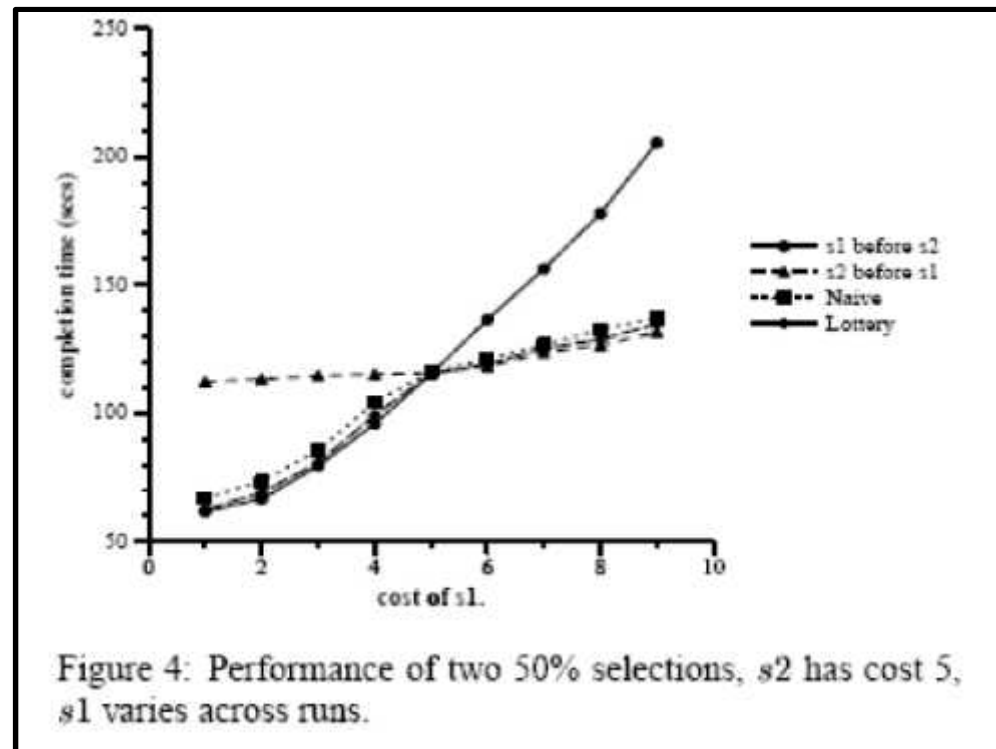
What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ Naïve approach naturally adjusts based on cost of operators
- ❖ Lottery also adjusts based on cost very well.

```
SELECT *  
FROM U  
WHERE s1() AND s2();
```



Experimental Results

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ Naïve approach does not adjust based on selectivity
- ❖ Lottery also adjusts based on selectivity




Figure 5: Performance of two selections of cost 5, s2 has 50% selectivity, s1 varies across runs.

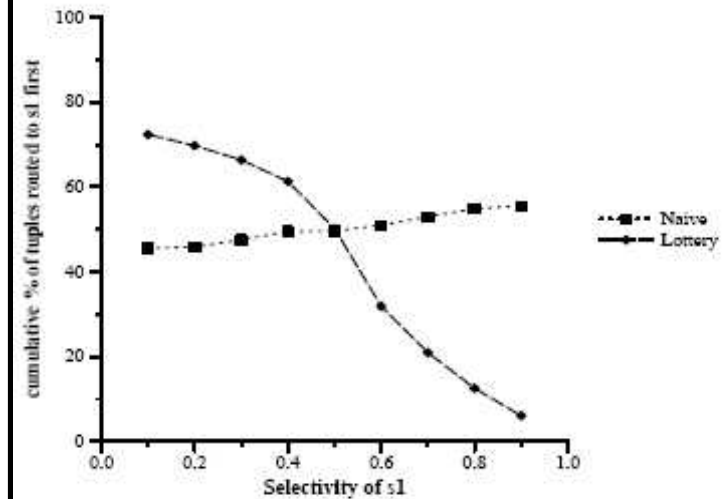


Figure 6: Tuple flow with lottery scheme for the variable-selectivity experiment (Figure 5).

Performance of Two Joins

Motivation

- Traditional query optimizers
- Problems

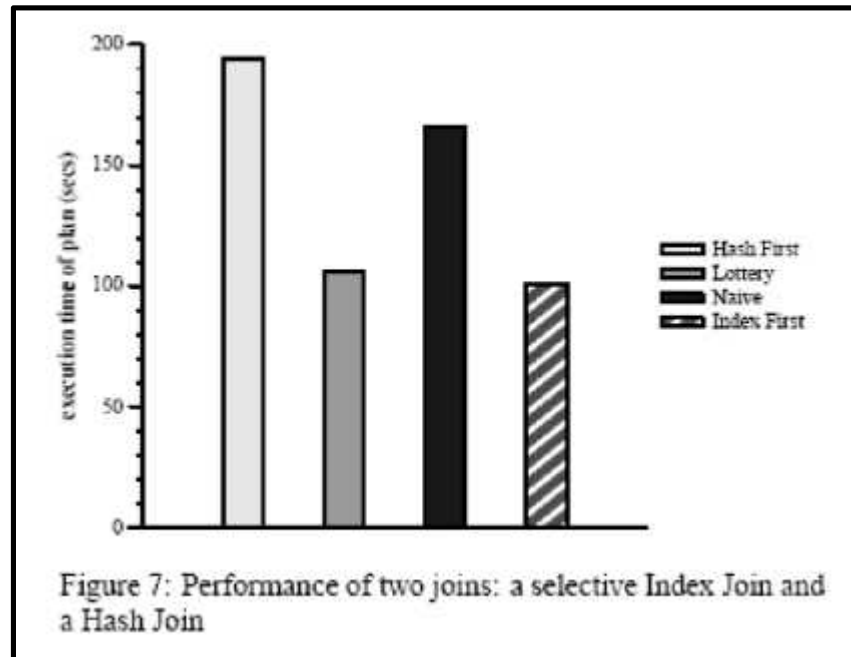
What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ Naïve perform between best and worst
- ❖ Lottery performs nearly optimal

```
SELECT *  
FROM R, S, T  
WHERE R.a = S.a  
AND S.b = T.b
```



Performance of Hash Joins

Motivation

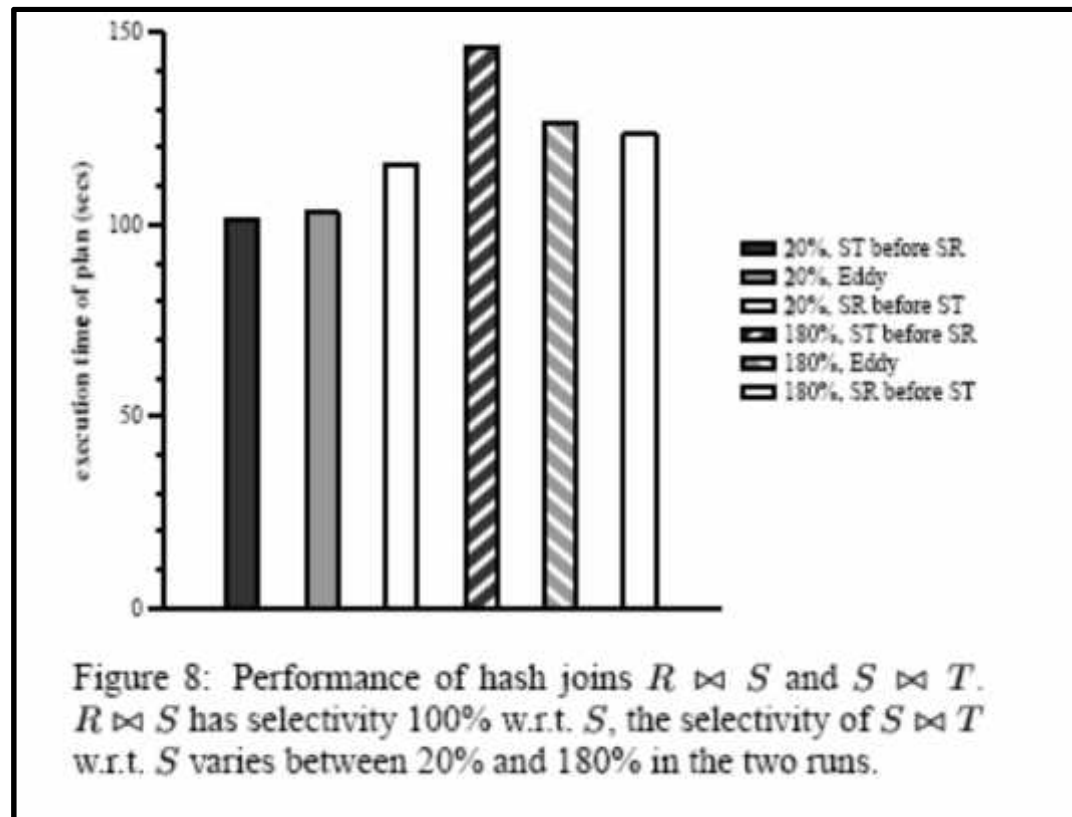
- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ All joins are ripple joins
 - ❖ Change selectivity of join predicate
 - ❖ In all cases eddy with lottery is close to optimal.



Performance of another Join

Motivation

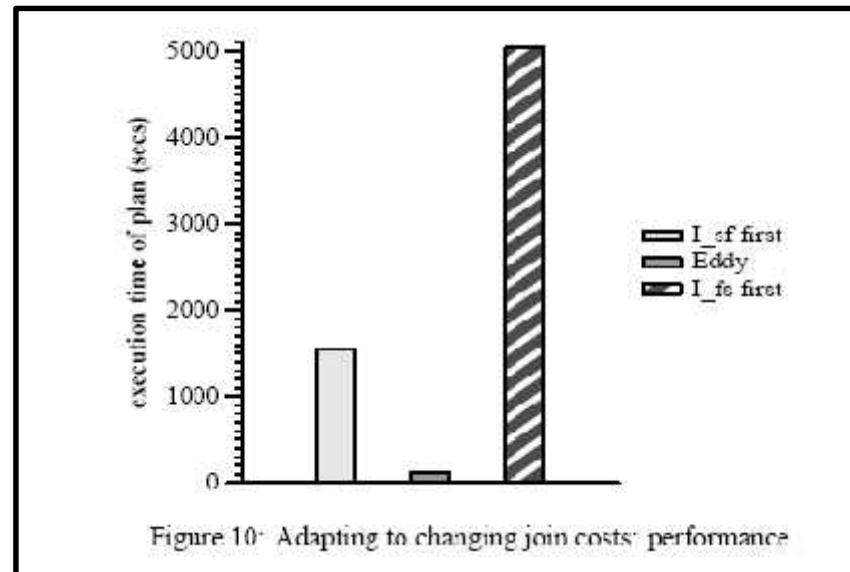
- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ 3-table equijoin query, two tables are external and used as “inner” relation by index join.
- ❖ Third relation has 30,000 tuples.
 - ❖ Eddy outdoes both static plans



Adaptability

Motivation

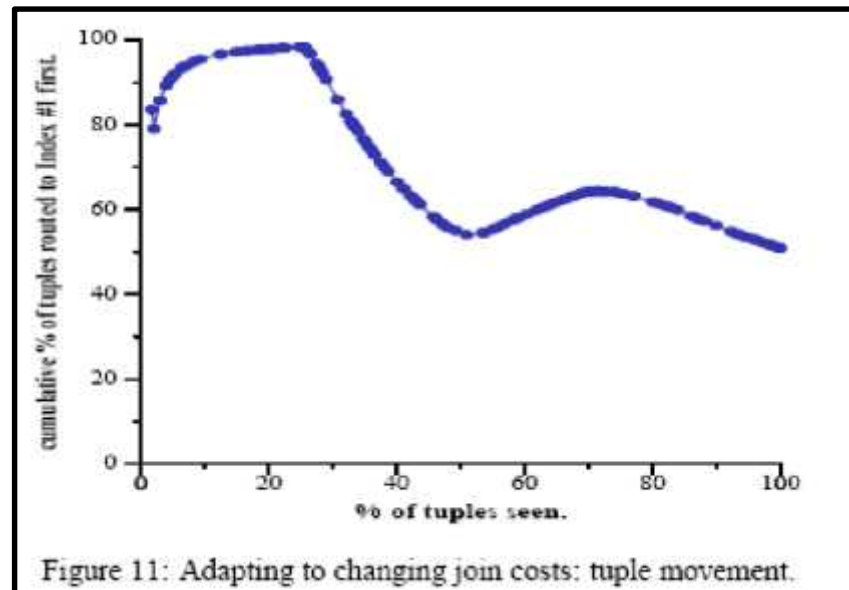
- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ Toggle the cost of operator 3 times throughout experiment
- ❖ Eddy switches how many tuples are first processed by that operator



Delayed Delivery

Motivation

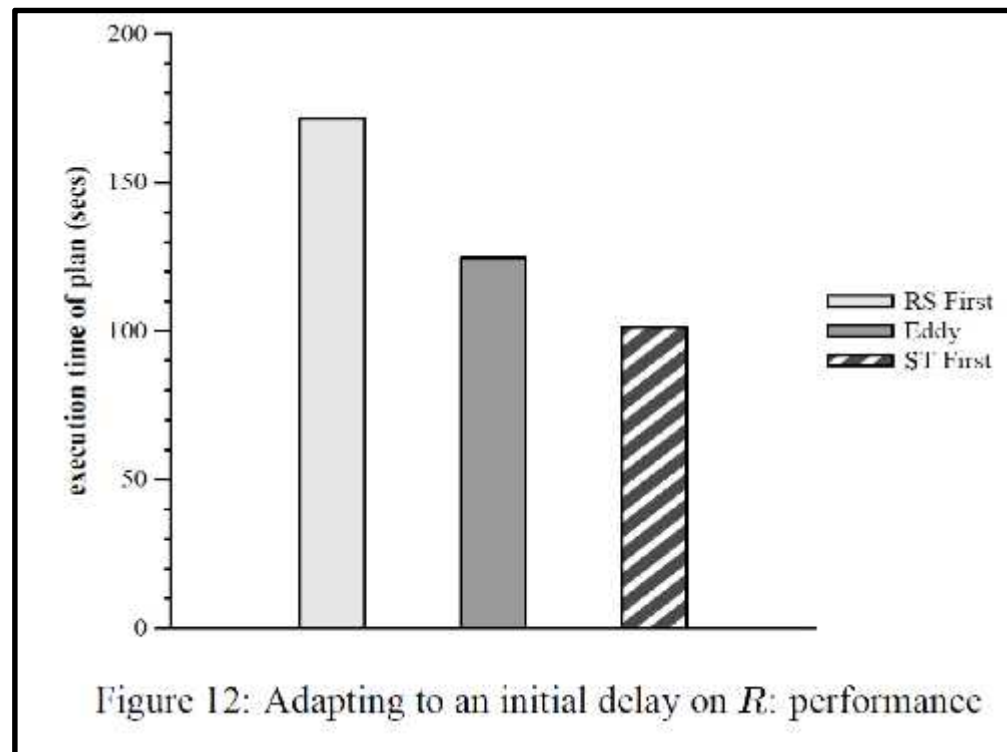
- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ RS Join does not produce any output tuples during early part of processing
 - ❖ Eddy awards most S tuples to RS joins initially
 - ❖ Ticket scheme does not capture the growing selectivity inherent in a join with delayed input.



Problems

Motivation

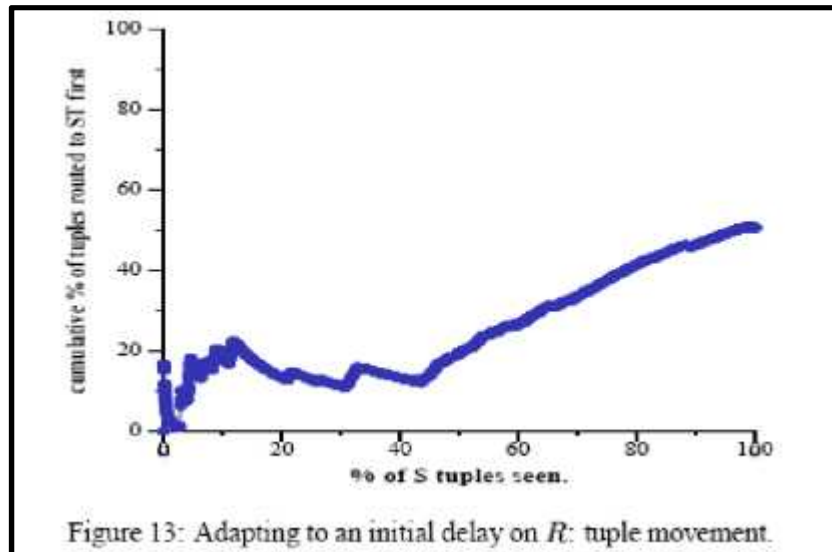
- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- **Experimental results**

Summary

- ❖ Does not work well when there is initially a long delay at an operator
- ❖ Eddy gives all tuples to operator because operator is not returning tuples that satisfy the join predicates criteria (not until after the delay)
- ❖ Eventually this problem is figured out by eddy (Just may take some time).



Re-optimization vs. Eddies

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Re-optimization (Kabra, Dewitt):
 - ❖ reordering queries at the end of pipelines
- ❖ Eddies:
 - ❖ Adaptively reorder pipelined operators on the fly
 - ❖ Learning algorithms that adaptively learns how to route the tuples to the pipelined operators

Strength of Eddy

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

- ❖ Per tuple adaptivity: Be beneficial for rapidly changing, unpredictable environments
- ❖ Can be used in concert with existing optimizers to improve adaptability within pipelines.

Summary

Motivation

- Traditional query optimizers
- Problems

What is eddy

- Reordering operators
- Eddies
- Routing tuples
- Experimental results

Summary

❖ Eddies are:

- ❖ A query processing mechanism that allow fine grained, adaptive, online optimization
- ❖ Beneficial in the unpredictable query processing environments

❖ Challenges:

- ❖ Develop ticket policies that can formally proved to converge quickly
- ❖ To attack the remaining static aspects
- ❖ To harness the parallelism and adaptivity available to us in rivers



Questions and Comments

Thank you very much!