


Granularity of Locks and **Degree of Consistency** in a Shared Data Base(PART 2)

Presented by: Jinglin Peng


Authors: J.N. Gray, R.A. Lorie, G.R. Putzolu, I.L. Traiger

You're designing a transaction system...

Requirement: account A wants to transfer \$100 to account B.



- ✓ Read balance from account A.
- ✓ Subtract \$100 from balance.
- ✓ Write new balance to account A.



- ✓ Read balance from account B.
- ✓ Add \$100 to balance.
- ✓ Write new balance to account B.

Timeline



Ideally, everything is alright. But how about in real world?

You're designing a transaction system...

Scenario 1:

Account A

- ✓ Read balance from account A.
- ✓ Subtract \$100 from balance.
- ✓ Write new balance to account A.



Power Failure

Account B

- × Read balance from account B.
- × Add \$100 to balance.
- × Write new balance to account B.

Timeline

Result: A lost \$100, but B didn't get the \$100

You're designing a transaction system...

Scenario 2:

Account A

- ✓ Read balance from account A.
- ✓ Subtract \$100 from balance.
- ✓ Write new balance to account A.

Account B

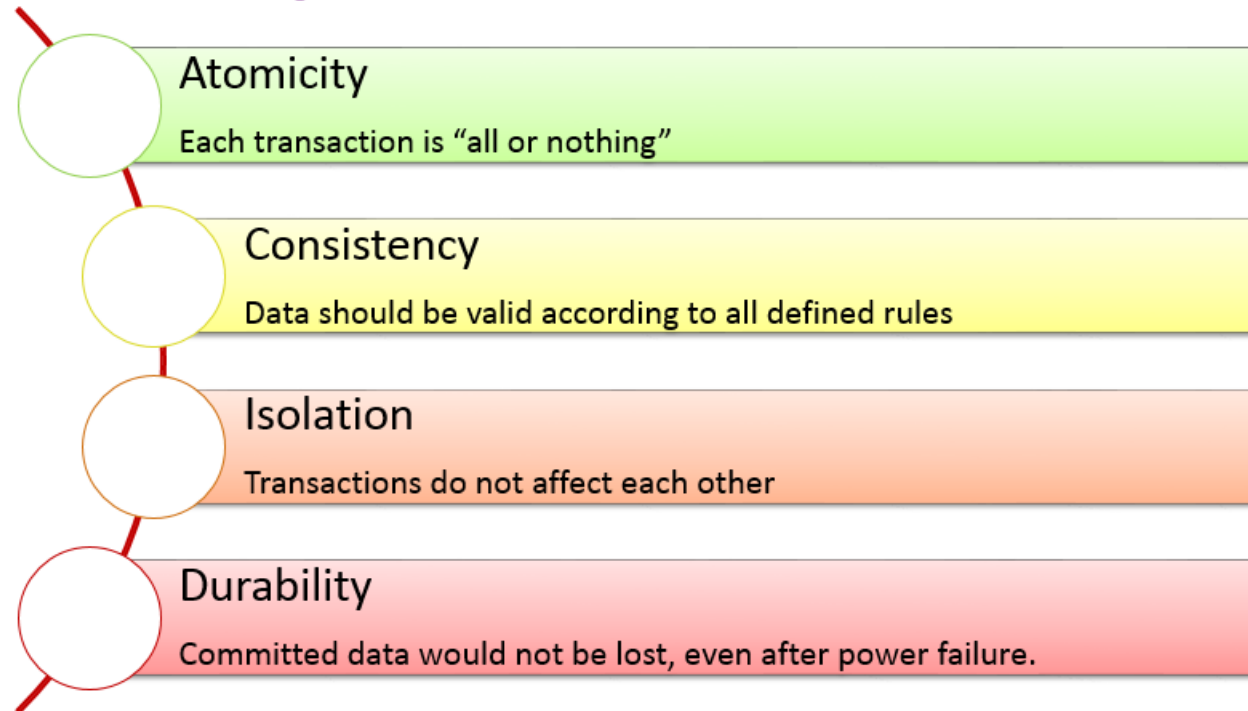
- ✓ Read **balance** from account B.
- × **B deposits \$50.**
- ✓ Add \$100 to **balance**.
- ✓ Write new balance to account B.

Timeline

Result: B lost the \$50 deposit


Requirement for Transaction System

ACID Properties



Today we'll talk about **consistency and isolation**.

Temporary Inconsistency

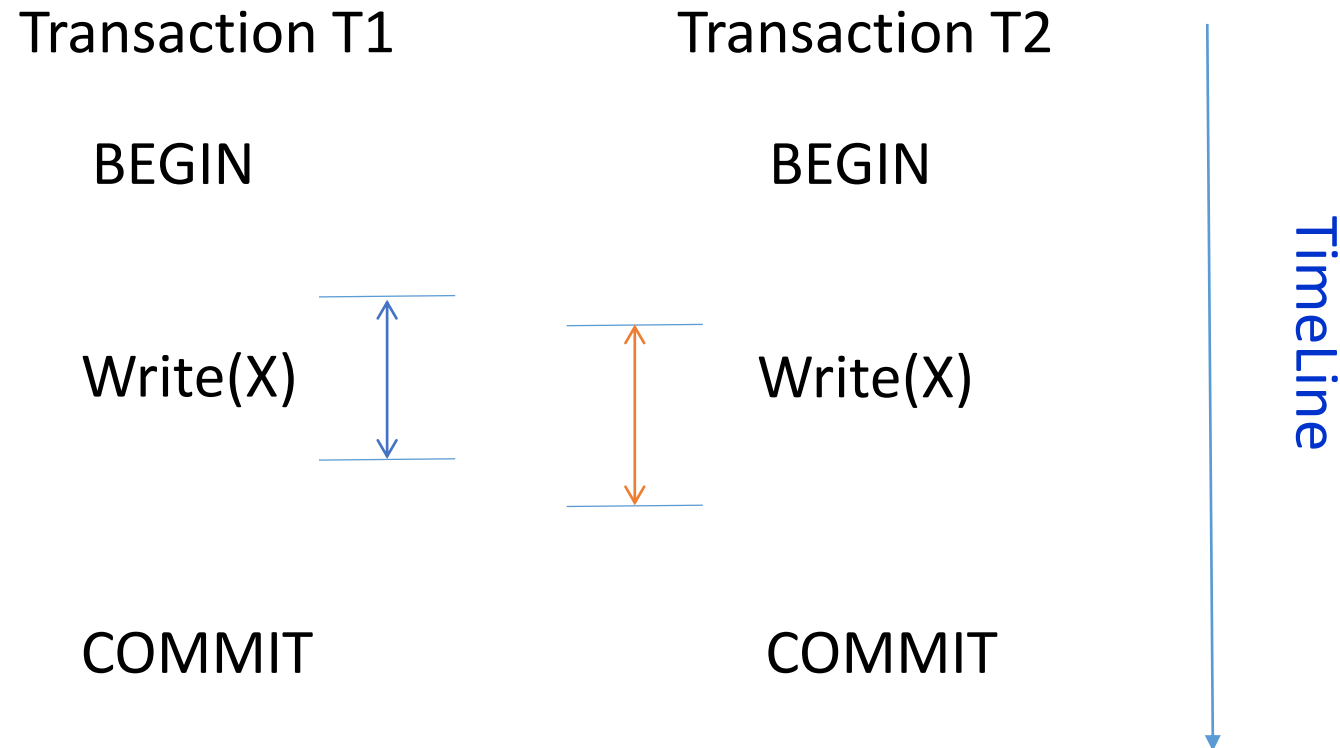


Among transaction, temporary inconsistency will occur.

Goal: keep consistency after transaction.

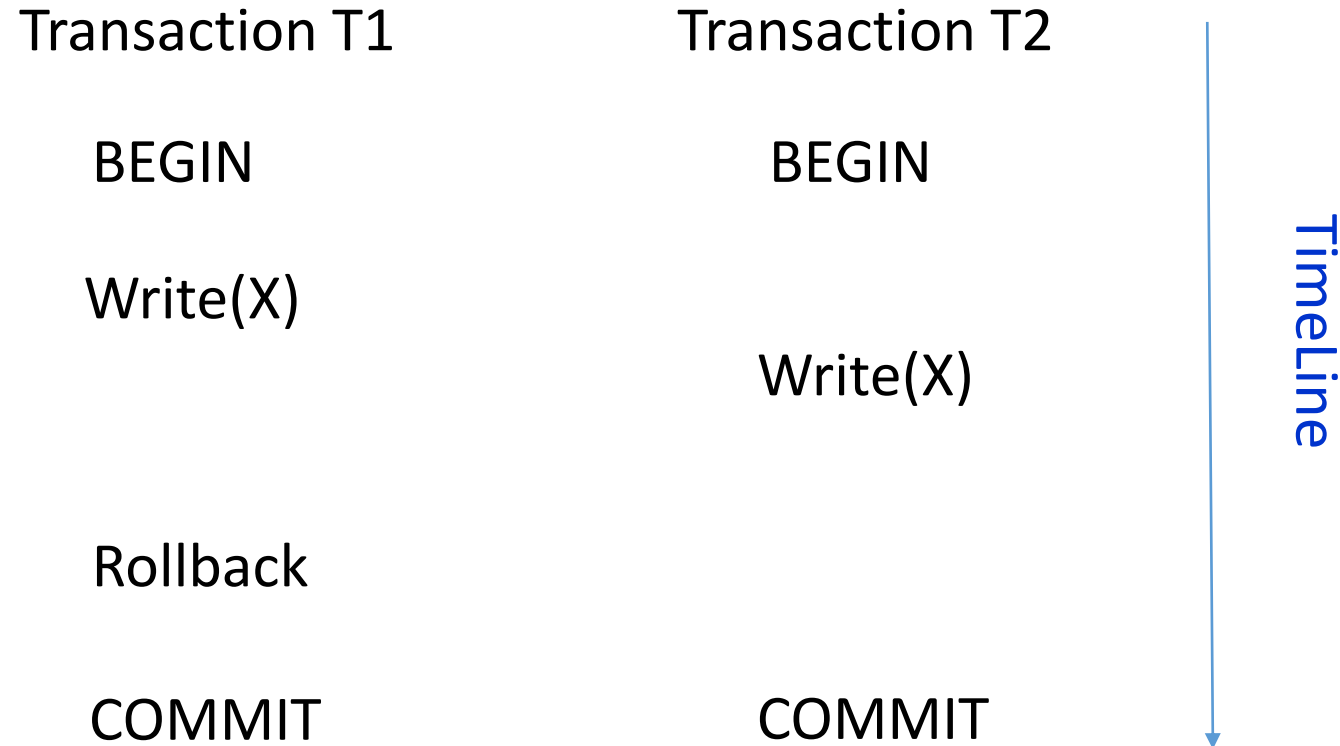
What's problem the temporary inconsistency may cause?

Problem 1: Garbage Read



Result: the value of X is unclear.

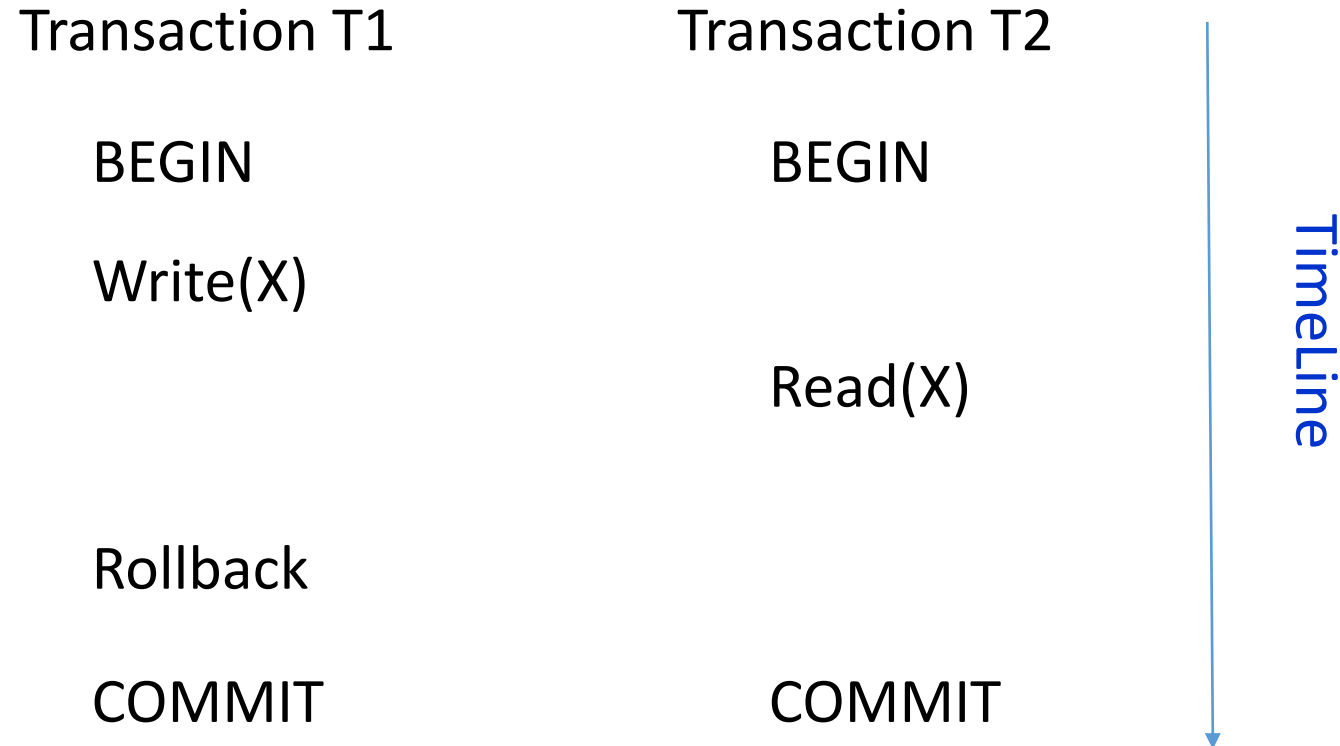
Problem 2: Lost Update



Result: update of transaction T2 is lost (rewritten by transaction T1)!

W-W conflict!

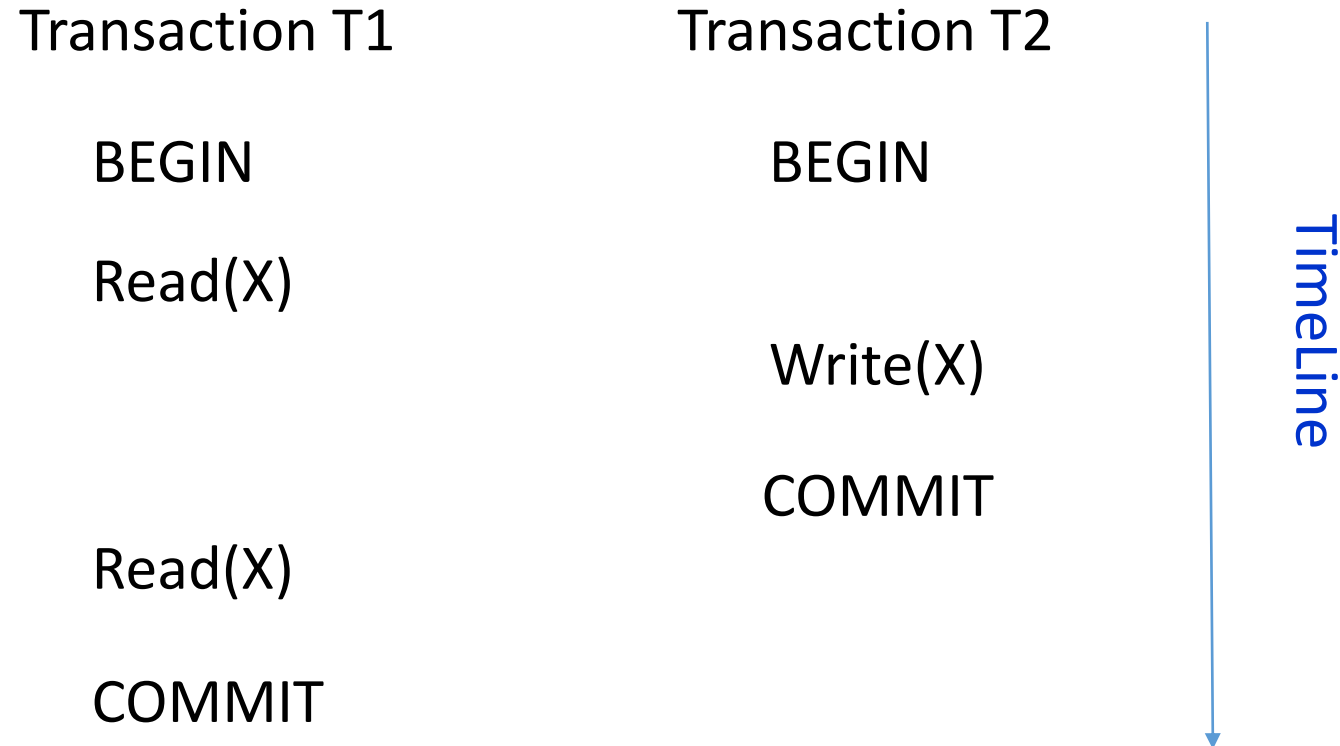
Problem 3: Dirty Read



Result: record X in transaction T2 is now dirty!

W-R conflict!

Problem 4: Unrepeatable Read



Result: transaction T1 might get a record with different values between reads!

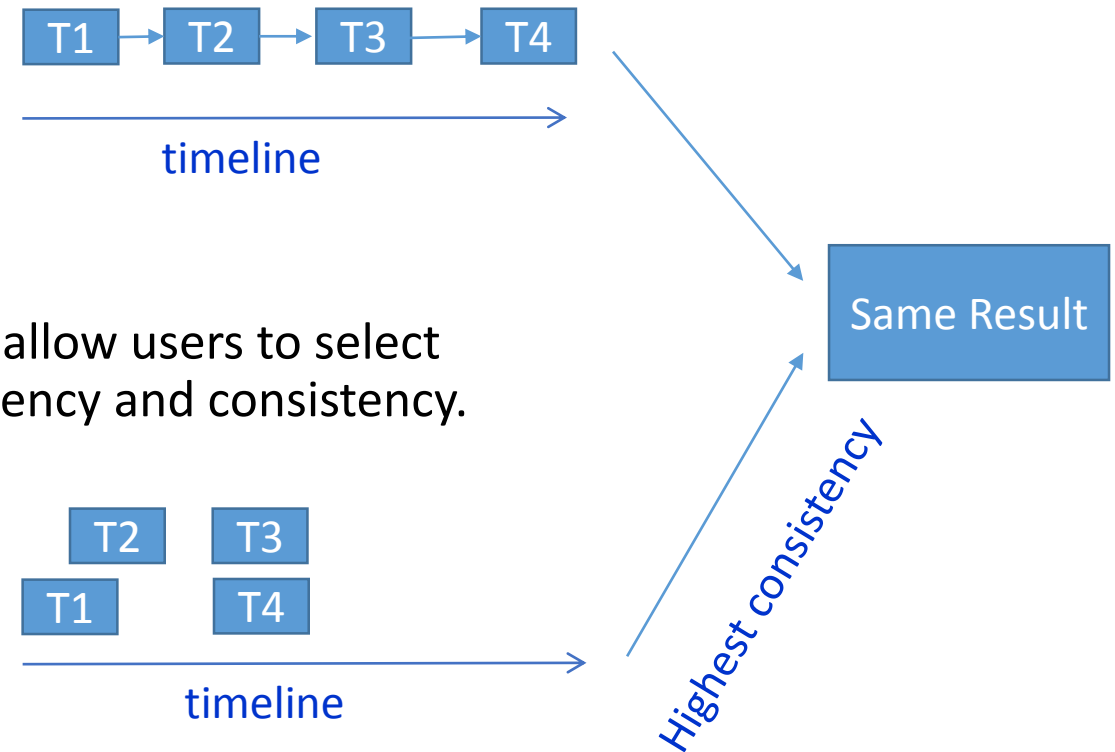
R-W conflict!

Prevention of Inconsistency

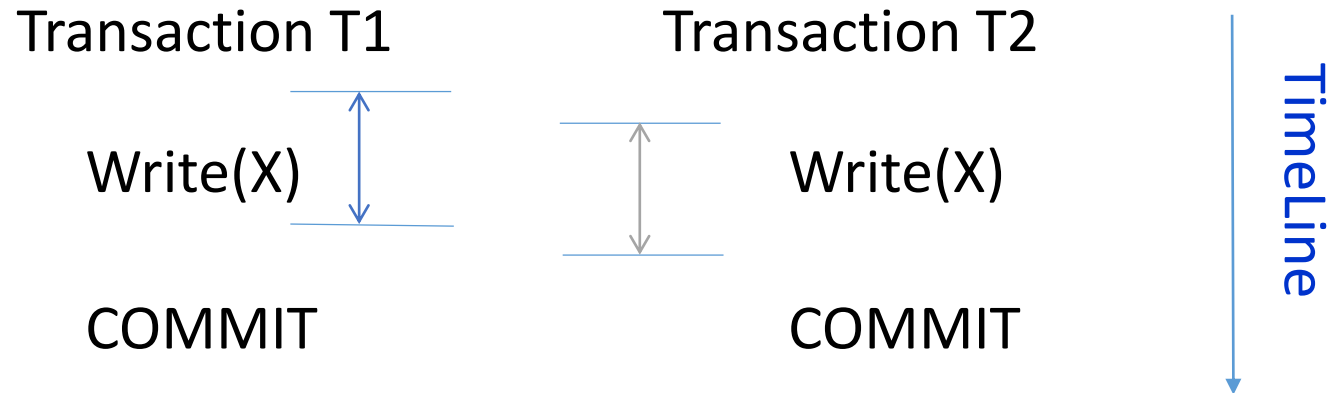
- Naïve Solution:
 - Execute transaction one at a time in a sequence.

- Drawback: latency!

- Smart Solution:
 - Classify consistency into different degrees, allow users to select degree to achieve the tradeoff between latency and consistency.



Prevention of Inconsistency: Degree 0

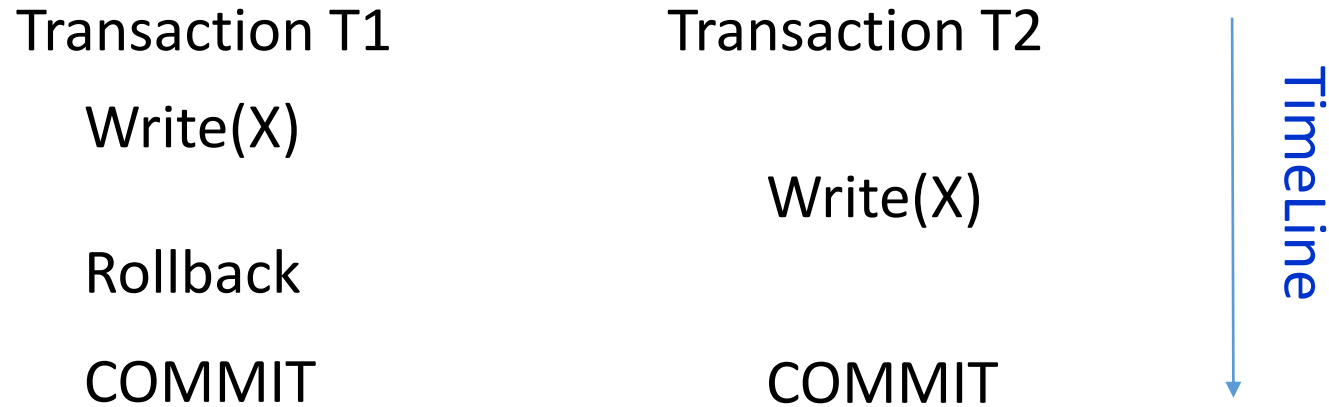


Problem: Who knows what value X will wind up holding? (garbage read)

Solution: set short write locks (short=until action finished)

- Before T1 write X, set a write lock to X
- After T1 write X, release the lock.

Prevention of Inconsistency: Degree 1

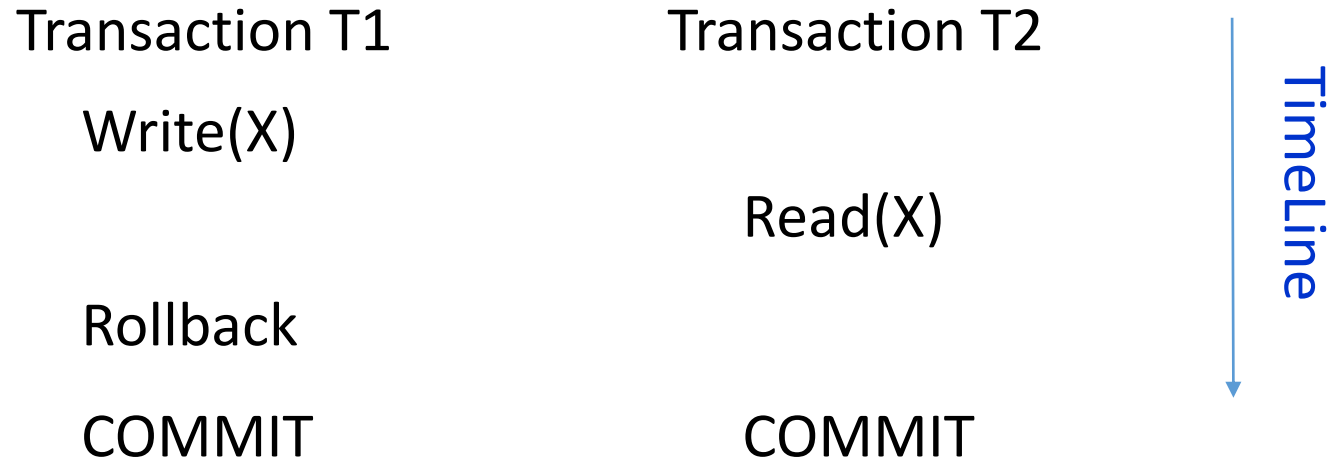


Problem: Update due to T2 is lost. (**lost update**)

Solution: set long write locks (long=until transaction committed)

- Before T1 write X, set a write lock to X
- After T1 is committed, release the write lock.

Prevention of Inconsistency: Degree 2




Problem: Now T2's read is bogus. (**dirty read**)

Solution: long write locks + short read locks

- Before T1 write X, set a write lock to X.
- Before T2 read, set a read lock to X.
- After finish T2's read, release the read lock.
- After T1 is committed, release the write lock.

Prevention of Inconsistency: Degree 3



Problem: Now T1 has two different values for X. (**unrepeatable read**)

Solution: long write locks + long read locks → 'two-phase locking'

- Before T1 read X, set a read lock to X.
- After T1 commit, release the read lock.
- Before T2 write X, set a write lock to X.
- After finish T2's write, release the write lock.

Conclusion of Consistency(Isolation) Degree

- Degree 0: short write locks -> garbage read
- Degree 1: long write locks -> lost update (W-W conflict)
- Degree 2: long write locks + short read locks -> dirty read (W-R conflict)
- Degree 3: long write locks + long read locks -> unrepeatable read (R-W conflict)

long = end of transaction

Short=end of action (write/read)

Definitions of Schedule

Problem: how to determine the consistency of a set transactions?

- Schedule:
 - A sequence of actions of a set of transactions.
- Schedule Consistency:
 - If **all** transactions run at degree 0 (1, 2 or 3) consistency in schedule S then S is said to be a degree 0 (1,2 or 3) consistent schedule.

Figure out how the transaction rely on each other.

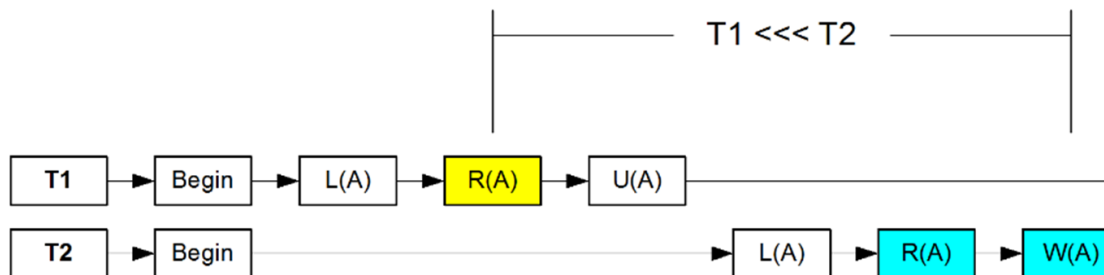
Transaction Interdependency

Suppose T1 & T2 perform action on same entity. T1 performs before T2.

Definition of Dependency

Dependency	Action	Degree
T1<T2	W->W	1
T1<<T2	W->W W->R	2
T1<<<T2	W->W W->R R->W	3

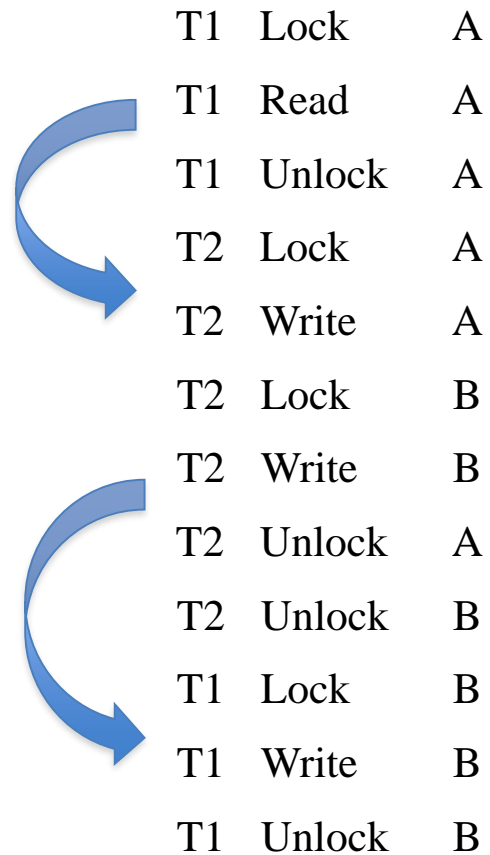
Example:



Assertion w.r.t Dependency

A schedule is degree 1 (2 or 3) consistent if and only if the closure of relation $<$ ($<<$ or $<<<$) is a partial order.

Example 1



Definition of Dependency

Dependency	Action	Degree
$T1 < T2$	$W \rightarrow W$	1
$T1 < < T2$	$W \rightarrow W \mid W \rightarrow R$	2
$T1 < < < T2$	$W \rightarrow W \mid W \rightarrow R \mid R \rightarrow W$	3


$T2 < T1, T2 < < T1, T2 < < < T1$

$T1 < < < T2$

$< < <$ is not partial order

- The schedule is degree 2 consistent but not degree 3 consistent.
 - T1 runs at degree 2 consistency, T2 runs at degree 3 consistency.

Example 2



Definition of Dependency

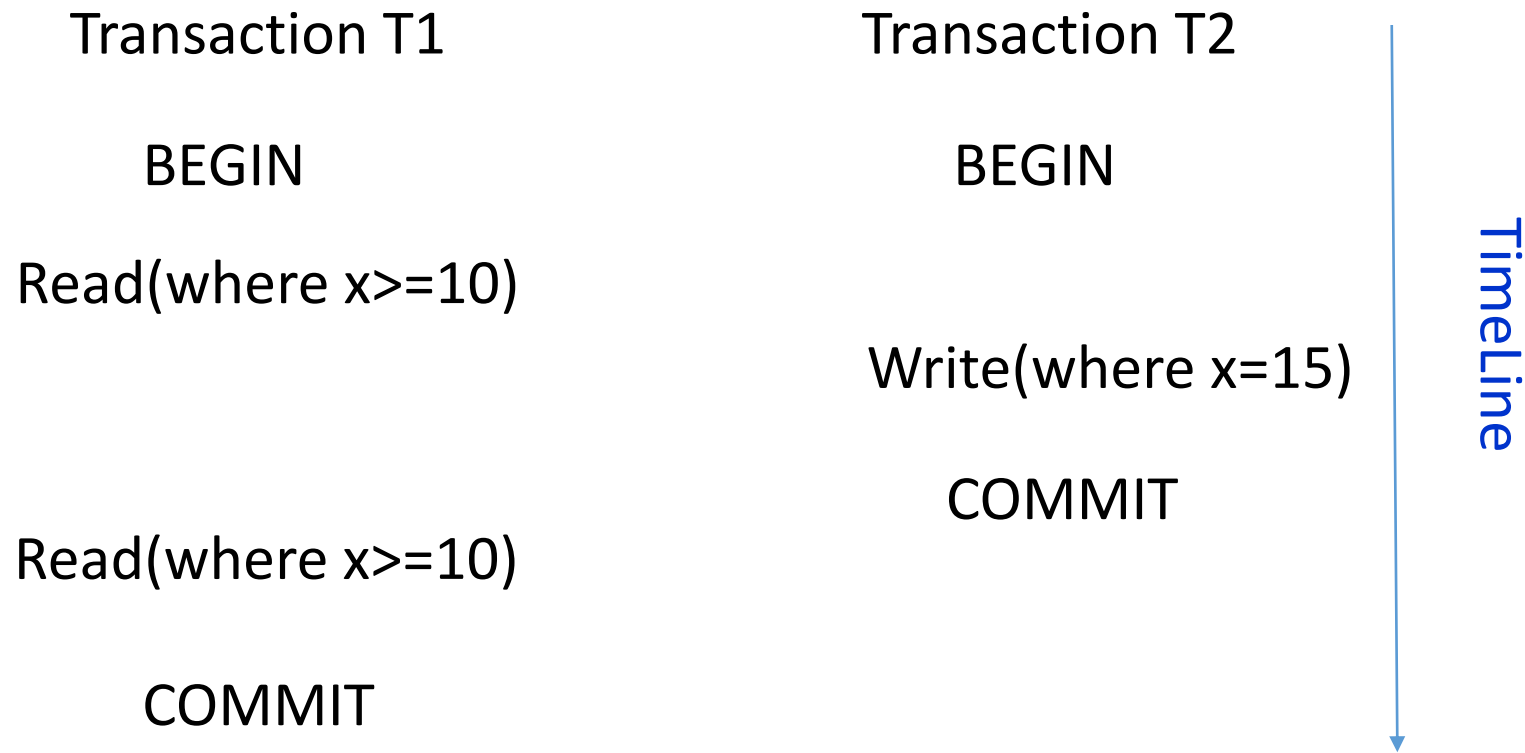
Dependency	Action	Degree
T1<T2	W->W	1
T1<<T2	W->W W->R	2
T1<<<T2	W->W W->R R->W	3

T2<T1, T2<<T1, T2<<<T1
T1<<T2, T1<<<T2

→ <<< and << are not partial order

- The schedule is degree 1 consistent.
 - T1 runs at degree 3 consistent, T2 runs at degree 1 consistent.

Problem out of the paper: Phantom Read



Results: results fetched by transaction T1 may be different in both reads.

Solution: range lock!

Summary

- 4 degrees of consistency: allow user tradeoff latency and consistency.
- Dependency among transactions: determine the consistency of schedule.

Thank you!

Q & A