

The Volcano Optimizer Generator: Extensibility and Efficient Search

Presenter: Zicun Cong

School of Computing Science
Simon Fraser University

CMPT843, February 6, 2016

Outline

Introduction


Method

Experiment


Conclusion

Query Execution

- ▶ **Logical Operator:** A function map the operator's inputs to its outputs. (e.g. join, selection, projection)
- ▶ **Physical Operator:** An algorithm that implements a logical operator. (e.g. hash join, merge join)
- ▶ **Operator Expression:** A hierarchy of operators.
- ▶ **Execution Plan:** An expression made up entirely of physical operators



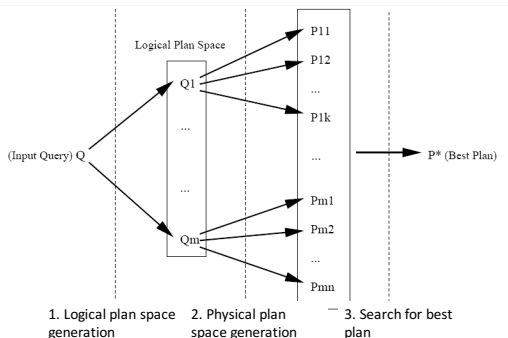
Logical Expression



Physical Expression

Query Optimization

- ▶ **Definition:** For a given query, find an execution plan for it that has the lowest cost
- ▶ **Challenges:** The number of candidate execution plans is huge
 - ▶ Equivalent expressions (e.g. $A \bowtie (B \bowtie C) = (A \bowtie B) \bowtie C$)
 - ▶ Different algorithms for each operation (e.g. Sort Merge Join, Hash Join, Nested Loop Join)



Motivation

Cost difference between evaluation plans for a query can be enormous (e.g. seconds vs. days)

Existing Methods:

- ▶ Heuristic search: Use static rules to generate plan
- ▶ Stratified search: Planning is done in multiple stages
- ▶ Unified search: Perform query planning all at once

Problems:

- ▶ **Not effective:** Hard to generate good plans for complex queries
- ▶ **Not extensible:** Rules maintenance is a huge pain
- ▶ **Not efficient:** Waste a lot of effort in searching

Overview of Volcano

Objectives

- ▶ Usability as a stand-alone tool
- ▶ More efficient resource usage
- ▶ Extensible support for physical properties

Optimizer Components

Physical property: The properties of physical layout of data

Logical Rules: Logic-Logic transformation

Implementation Rules: Logic-Implementation transformation

Algorithm/Enforcer: Have required physical properties of their inputs, physical properties of its output, and cost function

Example


- ▶ Physical property: Select * FROM A **ORDER BY id**
- ▶ Logic Rules: $S \bowtie R = R \bowtie S$
- ▶ Implementation Rules: $\bowtie \rightarrow \{\text{Merge Join, Hash Join, Sort + Hash Join}\}$
- ▶ Merge Join requires the inputs to be sorted on the join attributes and produces the output sorted

Optimizer Generator

Model Specification

- ▶ A set of logical operators and algebraic transformation rules
- ▶ A set of algorithms, enforcers and implementation rules
- ▶ An applicability function for each algorithm and enforcer
- ▶ A cost function for each algorithm and enforcer
- ▶ Property function for each operator, algorithm and enforcer

Generate Optimizer



Plan Search Engine

Basic Ideas

- ▶ The optimal execution plan of a query is composed of the optimal execution plan of its sub-queries.
- ▶ If the cost of a sub execution plan P is larger than cost threshold, any larger execution plan containing P can be pruned
- ▶ Recursively divide the query into sub-queries and find optimal execution plans of each sub-queries.

Example

SELECT * FROM A, B, C WHERE A.id = B.id AND B.id = C.id
 $\text{Cost}(A \bowtie (B \bowtie C)) = \text{Cost}(A) + \text{Cost}(\bowtie) + \text{Cost}(B \bowtie C)$

Plan Search Engine


Enumerate Search Space

Three possible moves

- ▶ Logic-Logic Transform
- ▶ Logic-Implementation Transform
- ▶ Enforcers for required physical property

Example

SELECT * FROM R, S, T WHERE R.x=S.x AND S.y=T.y
ORDER BY S.y



Algorithm Overview


```
FindBestPlan (LogExpr, PhysProp, Limit)
  /*Check look-up table (LogExpr, PhysProp) -> (Execution Plan, Cost)*/
  If (LogExpr, PhysProp) in table
    /*Prune high cost plans*/
    If cost < Limit return (Execution Plan, Cost)
    else return Fail
  /* else: optimization required */
  Generate possible moves
  For move in moves
    Handle the move
  /* maintain the look-up table of explored facts */
  If LogExpr is not in the look-up table:
    Update loop-up table
  return best Plan and Cost
```

Handle Moves

Handle Logical Transformation

Recursively call the **FindBestPlan** function with the new logical expression

Examples




Handle Moves

Handle Implementation Transformation

1. Update the total cost. TotalCost = Algorithm Cost
2. Recursively call FindBestPlan on each input of the algorithm

Examples




Handle Moves

Handle Enforcer Transformation

1. Update physical property and cost
2. Recursively call the **FindBestPlan** function with the new physical property and cost

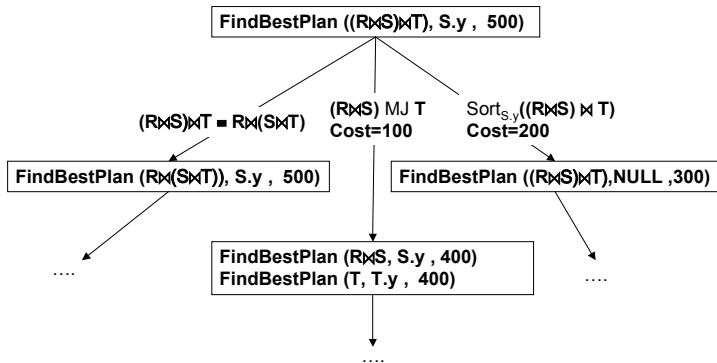
Examples




Depth-First Search Tree

Query

SELECT * FROM R, S, T WHERE R.x=S.x AND S.y=T.y
ORDER BY S.y



Experiment



Discussion

Pros

- ▶ Unified, fully cost-based model
- ▶ Easily add new operations and equivalence rules
- ▶ Pruning strategies can effectively reduce the searching space

Cons

- ▶ Only compare their method with one query optimization method
- ▶ Exhaustively generates all logically equivalent expressions

Conclusion

- ▶ Volcano improves the work of EXODUS
- ▶ Top-down optimizer, uses dynamic programming and memoisation
- ▶ Enumerates physical search space in depth-first order
- ▶ Uses branch and bound pruning to prune the search space
- ▶ Use physical properties to direct the search
- ▶ Volcano has gained widespread acceptance in the industry as a state-of-the-art optimizer; the optimizers of Microsoft SQL Server [21] and Tandem ServerWare SQL Product [6] are based on Volcano

Backup Slids

Logical / Implementation Rules

- ▶ **Logical Rules:** The algebraic rules of expression equivalence. They are used to get equivalent logical expressions.
- ▶ **Implementation Rules:** The possible mappings of operators to algorithms
- ▶ Algorithms have three properties
 - ▶ Required physical properties of its inputs
 - ▶ Physical properties of its output
 - ▶ Cost function

Example

- ▶ Commutative law: $R \bowtie S = S \bowtie R$
- ▶ Join can be implemented by merge-join, hash join or nested-loop join.
- ▶ merge-join requires the inputs to be sorted on the join attributes and produces the output sorted

Optimizer Generator

Model Specification

- ▶ A set of logical operators and algebraic transformation rules
- ▶ A set of algorithms, enforcers and implementation rules
- ▶ An applicability function for each algorithm and enforcer
- ▶ A cost function for each algorithm and enforcer
- ▶ Property function for each operator, algorithm and enforcer

- ▶ Given a model specification, Volcano generate a query optimizer.
- ▶ Optimizer source code is compiled and linked with the other DBMS software such as the query execution engine and with the search engine.
- ▶ When the DBMS is operational, the input query is passed to the optimizer, which generates an optimized plan for it.

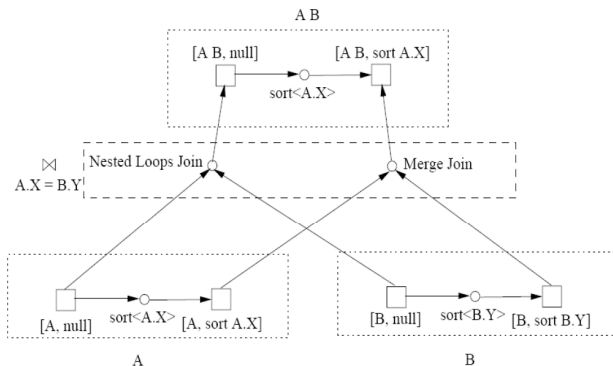
Difference between EXODUOS

- ▶ Volcano took less time to optimize
- ▶ EXODUS optimizer generator measurements were quite volatile and took a lot of memory
- ▶ EXODUSs generated optimizer and search engine do not explore and exploit physical properties

Difference between Starburst

- ▶ Volcano took less time to optimize
- ▶ New operators are integrated at the query rewrite level
optimization in this level is heuristic
- ▶ Starburst has a hierarchy of intermediate levels

Physical Transformation



PQDAG for $A \bowtie B$ on $A.X=B.Y$

dotted boxes - logical eqv. nodes; solid boxes - phys. eqv. nodes

Algorithm

FindBestPlan (LogExpr, PhysProp, Limit)

- if the pair LogExpr and PhysProp is in the look-up table
 - if the cost in the look-up table < Limit
 - return Plan and Cost
 - else return failure
- /* else: optimization required */
- create the set of possible "moves" from
 - applicable transformations
 - algorithms that give the required PhysProp
 - enforcers for required PhysProp
- order the set of moves by promise

Algorithm

- for the most promising moves
 - if the move uses a transformation
 - apply the transformation creating NewLogExpr
 - call FindBestPlan (NewLogExpr, PhysProp, Limit)
 - else if the move uses an algorithm
 - TotalCost := cost of the algorithm
 - for each input I while TotalCost < Limit
 - determine required physical properties PP for I
 - Cost = FindBestPlan (I, PP, Limit – TotalCost)
 - add Cost to TotalCost
 - else /* move uses an enforcer */
 - TotalCost := cost of the enforcer
 - modify PhysProp for enforced property
 - call FindBestPlan for LogExpr with new PhysProp

Algorithm

- /* maintain the look-up table of explored facts */
- if LogExpr is not in the look-up table
 - insert LogExpr into the look-up table
 - insert PhysProp and best plan found into look-up table
- return best Plan and Cost